Cloud Provider Advisor

Product Requirement Document — Combined Edition

On-Premise → Cloud Migration Decision Assistant

Version 2.0 | March 2026

About This Document

This document merges two independent specification drafts — one authored via ChatGPT and one authored via Claude — into a single, authoritative Product Requirement Document. The best elements of each have been preserved, reconciled, and extended. The ChatGPT draft contributed a clean decision algorithm structure, a concise 25-question bank, and a clear UX screen-by-screen flow. The Claude draft contributed a comprehensive compliance and security layer, extended provider coverage (IBM Cloud, Oracle Cloud), a weighted scoring rationale framework, hard-eliminator logic, and detailed acceptance criteria and KPIs. Together they form the most complete specification available for this feature.

1. Overview & Problem Statement

1.1 Purpose

The Cloud Provider Advisor is an intelligent, questionnaire-driven feature that guides users transitioning from on-premise or local storage environments to cloud infrastructure. By collecting structured inputs across workload type, technical stack, compliance posture, and budget, the system produces a ranked recommendation of the most suitable cloud provider — acting as a trusted bridge through what is otherwise an overwhelming decision.

1.2 Problem Statement

Startups and organizations moving off local infrastructure commonly face:

- No familiarity with AWS, Azure, GCP, IBM Cloud, or Oracle Cloud differences
- Confusion over pricing models — pay-as-you-go vs. reserved vs. spot pricing
- Uncertainty about which compliance certifications (HIPAA, GDPR, FedRAMP, PCI-DSS) each provider satisfies
- Risk of choosing a provider that is difficult to migrate away from (vendor lock-in)
- Analysis paralysis from hundreds of competing services across providers
- Delayed migration decisions resulting in continued on-premise cost drag

A structured recommendation engine replaces guesswork with a repeatable, explainable decision process.

1.3 Goals

Business Goals

- Improve onboarding experience for cloud migration users
- Position the product as a cloud advisory and migration planning tool
- Collect infrastructure requirement data to power future cost-estimation features
- Increase platform engagement and reduce time-to-first-migration

User Goals

- Quickly identify the best-fit cloud provider in under 8 minutes
- Understand clearly why a specific provider was recommended
- Reduce research time and avoid costly post-migration re-platforming
- Gain confidence before committing to a cloud vendor

Non-Goals (v1)

- Automatic cloud account provisioning or migration execution
- Live cost simulation with real-time pricing API calls (Phase 2)
- Production-ready architecture generation (Phase 3)

1.4 Target Users

Persona	Role	Primary Need
Startup Founder	Technical / Business	Simple, fast recommendation without cloud expertise
IT Manager	CTO / IT Director	Cost-effective, compliant platform for the organization
Solutions Architect	Cloud / Infra Architect	Best-fit services for a specific technical architecture
DevOps Engineer	Platform Engineer	Cloud-native tooling and CI/CD pipeline compatibility

Compliance Officer	Legal / Security	Regulatory alignment and data residency guarantees

2. Cloud Provider Decision Algorithm

2.1 Evaluation Dimensions & Weights

Each answer contributes a delta score to one or more providers, multiplied by the category weight. The final score is normalized to 100.

Formula	Final Score = Σ (Answer Delta × Category Multiplier) / Max Possible Score × 100

Category	Multiplier	Rationale
Compliance & Regulatory Requirements	**×3.0**	Legal mandates are non-negotiable hard constraints
Primary Workload Type	**×2.5**	Technical fit is the core utility dimension
Hybrid / Integration Needs	**×2.5**	Architecture compatibility is often binary
Data Volume & Storage	**×2.0**	Egress costs are the most common hidden expense
Budget & Cost Model	**×2.0**	Constrains the entire feasible solution space
Security & Encryption	**×2.0**	Often driven by the same regulatory mandates
AI / ML Requirements	**×1.5**	High differentiator for a subset of users
Support Level	**×1.5**	Critical for enterprises, lower priority for startups
Vendor Ecosystem Fit	**×1.5**	Existing tooling creates real migration switching costs
Scalability Needs	**×1.5**	Affects architecture choices from day one
Open-Source Preference	**×1.0**	Strong preference, rarely a sole decision-blocker
Existing Vendor Relationship	**×0.5**	Noted but kept low-weight for objectivity

2.2 Hard Eliminators (Pre-Scoring Filters)

Before weighted scoring begins, the following conditions immediately remove incompatible providers:

- FedRAMP / GovCloud required → eliminates any provider without an approved GovCloud region

- Data must reside in country X → eliminates providers with no data centre in that country
- HIPAA required → eliminates providers that have not signed a Business Associate Agreement (BAA)
- Budget < $500/month → downranks enterprise-tier-only providers
- Active-Active DR required → downranks providers with fewer than 3 global regions
- Strict BYOK / HSM required → eliminates providers without customer-managed key support

2.3 Score Initialization

Initialize	AWS = 0 | Azure = 0 | GCP = 0 | IBM = 0 | Oracle = 0

2.4 Score Deltas by Workload & Stack

Workload Type

Workload	AWS	GCP	Azure	IBM / Oracle
Web Applications / APIs	AWS +3	GCP +2	Azure +1	—
Data Analytics / Warehousing	AWS +2	GCP +3	Azure +1	—
AI / Machine Learning	AWS +2	GCP +3	Azure +2	—
Enterprise Applications	AWS +2	Azure +3	GCP +1	IBM +3
Containers / Kubernetes	AWS +2	GCP +3	Azure +2	—
Oracle DB / ERP Workloads	—	—	—	Oracle +4
High-Performance Computing	AWS +3	GCP +2	Azure +2	—

Technical Stack

Tech Stack	AWS	GCP	Azure	IBM/Oracle
.NET / Microsoft Stack	Azure +4	AWS +1	GCP +0	—
Open Source / Linux	AWS +3	GCP +3	Azure +1	—
Kubernetes / Containers	GCP +3	AWS +2	Azure +2	—
Oracle / Java EE	Oracle+4	IBM +2	AWS +1	—
Python / Data Science	GCP +3	AWS +2	Azure +1	—
IBM / AIX / Db2	IBM +4	AWS +1	—	—

Budget Sensitivity

Budget Tier	AWS	GCP	Azure
Highly cost-sensitive	AWS +2	GCP +3	Azure +1
Moderate sensitivity	AWS +2	GCP +2	Azure +2
Budget not a concern	AWS +2	GCP +1	Azure +3

Scalability & Cloud Expertise

Dimension	AWS	GCP	Azure
Rapid scaling expected	AWS +3	GCP +2	Azure +2
Moderate growth	AWS +2	GCP +2	Azure +2
No cloud experience	AWS +1	Azure +2	GCP +2
Experienced DevOps team	AWS +3	GCP +2	Azure +2

2.5 Recommendation Output

Provider	Final Score	Recommendation
AWS	78 / 100	**RECOMMENDED**
GCP	65 / 100	Alternative
Azure	52 / 100	—

Logic	Provider = MAX(Σ weighted scores). In a tie, the provider with higher scores on the user's highest-weight categories wins.

3. Cloud Provider Capability Matrix

The matrix below scores each provider (1–10) across eight capability dimensions. This informs both the scoring engine and the explanation text shown to the user on the recommendation screen.

Capability	Weight	AWS	Azure	GCP
Scalability	**20%**	9 ★★★★★	8 ★★★★☆	8 ★★★★☆
Cost Efficiency	**15%**	7 ★★★☆☆	7 ★★★☆☆	9 ★★★★★
Developer Ecosystem	**15%**	10 ★★★★★	8 ★★★★☆	8 ★★★★☆
AI / Machine Learning	**10%**	8 ★★★★☆	8 ★★★★☆	10 ★★★★★
Enterprise Integration	**10%**	7 ★★★☆☆	10 ★★★★★	7 ★★★☆☆

Containers / Kubernetes	**10%**	8 ★★★★☆	8 ★★★★☆	10 ★★★★★
Migration Tooling	**10%**	9 ★★★★★	9 ★★★★★	7 ★★★☆☆
Ease of Use	**10%**	7 ★★★☆☆	8 ★★★★☆	9 ★★★★★

3.1 Provider Best-Fit Profiles

Amazon Web Services (AWS)

Best for web apps, SaaS platforms, and polyglot enterprise architectures needing the broadest service catalogue and global reach. Strongest migration tooling (AWS Migration Hub, DMS, Server Migration Service).

- Watch for: Complex pricing structure; egress fees can surprise cost-sensitive teams

Microsoft Azure

Best for organizations on the Microsoft stack (Active Directory, .NET, SQL Server, Office 365) or those needing the strongest hybrid cloud connectivity via Azure Arc.

- Watch for: Portal complexity; premium support tiers can be expensive

Google Cloud Platform (GCP)

Best for data-intensive workloads, AI/ML products, and teams that value open-source standards. Kubernetes originated here; BigQuery and Vertex AI are best-in-class.

- Watch for: Smaller partner network; fewer regions in some geographies vs. AWS/Azure

IBM Cloud

Best for enterprises with existing IBM software (Db2, WebSphere, MQ) and financial services firms needing LinuxONE or Z integration. Strong regulated-industry compliance story.

- Watch for: Narrower general-purpose catalogue; smaller developer community

Oracle Cloud Infrastructure (OCI)

Best for Oracle Database and ERP workloads under the Bring Your Own License programme. Often the lowest-cost option for Oracle-licensed software.

- Watch for: Limited appeal outside the Oracle ecosystem; smaller ISV partner network

4. Expert Assessment Questionnaire (25 Questions)

Questions are grouped into five thematic stages. Smart branching reduces the presented questions to 12–15 for most users. Each card shows the scoring weight and a branching logic note for the engineering team.

Stage A — Infrastructure & Data (Q1–Q5)

Q1 How much total data do you currently store on-premise?

- Less than 1 TB
- 1 TB – 10 TB
- 10 TB – 100 TB
- 100 TB – 1 PB
- More than 1 PB

Weight: High **Branching note:** *Drives egress cost sensitivity and storage class recommendations across all providers.*

Q2 What percentage of your stored data changes or is written daily?

- Less than 5% — mostly read-only archive
- 5–30% — moderate transactional activity
- More than 30% — highly active datasets

Weight: Medium **Branching note:** *High churn favours object storage with low PUT pricing; skip Q3 if >30%.*

Q3 Do you require archival or cold storage for infrequently accessed data?

- Yes — significant volume of cold/archival data
- No — all data is actively accessed

Weight: Medium **Branching note:** *Skip if previous answer was >30% daily change.*

Q4 What is the largest single file or object you need to store?

- Under 5 GB
- 5 GB – 500 GB
- Over 500 GB (e.g., video, large DB dumps, genomics)

Weight: Low **Branching note:** *Very large objects require multipart upload support — all major providers support this.*

Q5 Do you require your application to be available across multiple geographic regions simultaneously?

- Yes — active-active multi-region (simultaneous)
- Yes — active-passive (failover only)
- No — single region is sufficient

Weight: High **Branching note:** *Active-active multi-region triggers the DR hard-eliminator filter.*

Stage B — Application Architecture (Q6–Q10)

Q6 Is your application monolithic or microservices-based?

- ► Monolithic — single deployable unit
- ► Microservices — independently deployable services
- ► Mixed / transitioning to microservices

Weight: Medium **Branching note:** *Microservices boost container/Kubernetes scoring for GCP and AWS.*

Q7 Do you currently use or plan to use containers (Docker / Kubernetes)?

- ► Yes — actively using in production
- ► Yes — planning to adopt during migration
- ► No — no container plans

Weight: High **Branching note:** *Yes answers add +3 GCP, +2 AWS, +2 Azure.*

Q8 What type of database does your application use?

- ► Relational (PostgreSQL, MySQL, SQL Server, Oracle DB)
- ► NoSQL (MongoDB, Cassandra, DynamoDB-style)
- ► Both relational and NoSQL
- ► Data warehouse / analytical DB (Redshift, BigQuery style)
- ► Oracle DB specifically

Weight: High **Branching note:** *Oracle DB answer triggers Oracle Cloud scoring boost.*

Q9 Is your application stateless or stateful?

- ► Stateless — any instance is interchangeable
- ► Stateful — sessions or data tied to specific instances

Weight: Medium **Branching note:** *Stateful applications require persistent storage and careful migration planning.*

Q10 Do you require automatic scaling based on traffic load?

- ► Yes — auto-scaling is essential to our architecture
- ► Maybe — we expect growth but haven't implemented it yet
- ► No — fixed capacity is sufficient

Weight: High **Branching note:** *Auto-scaling requirement adds +3 AWS, +2 GCP, +2 Azure.*

Stage C — Performance & Compliance (Q11–Q17)

Q11 What is your expected number of monthly active users at launch?

- ► Under 10,000
- ► 10,000 – 1 million
- ► Over 1 million

Weight: Medium **Branching note:** *High MAU drives availability and CDN requirements.*

Q12 Do you experience or expect significant traffic spikes (e.g., flash sales, seasonal peaks)?

- ► Yes — spikes can be 10× or more of baseline
- ► Moderate — occasional peaks of 2–5×
- ► No — traffic is relatively steady

Weight: Medium **Branching note:** *Spikes boost rapid-scaling providers (AWS +3, GCP +2).*

Q13 What are your application latency requirements for end users?

- ► Under 50 ms (real-time / gaming / trading)
- ► 50–200 ms (standard web application)
- ► 200 ms+ acceptable (batch / async workloads)

Weight: Medium **Branching note:** *Sub-50ms requirements require edge/CDN layer; all major providers support this.*

Q14 Which regulatory compliance frameworks must your cloud environment satisfy?

- ► HIPAA (health data — USA)
- ► GDPR (personal data — EU)
- ► PCI-DSS (payment card industry)
- ► FedRAMP / GovCloud (US Government)
- ► ISO 27001 / SOC 2 (general enterprise)
- ► ITAR / EAR (defence / export control)
- ► None required
- ► Multiple of the above

Weight: HIGH (×3.0) **Branching note:** *Hard-eliminator logic runs on this answer before any scoring.*

Q15 Where must your data physically reside?

- ► Specific country (e.g., must stay within India / EU / Australia)
- ► Specific continent only
- ► No restriction — any global region is acceptable

Weight: HIGH (×3.0) **Branching note:** *Country-specific residency eliminates providers without a local data centre.*

Q16 How sensitive is the data you will store or process?

- ► Public / non-sensitive only
- ► Internal business data (mildly sensitive)
- ► Confidential customer PII
- ► Highly classified (financial, healthcare, government records)

Weight: High **Branching note:** *PII and classified data boost encryption scoring weight.*

Q17 Do you require a specific disaster recovery strategy?

- ► Basic backup and restore (hours of recovery time acceptable)
- ► Warm standby (minutes to recovery)
- ► Active-Active multi-region (near-zero RTO)
- ► No formal DR plan yet

Weight: High **Branching note:** *Active-Active DR triggers the multi-region hard-eliminator filter.*

Stage D — Development Ecosystem (Q18–Q21)

Q18 What is your team's primary programming language or technology stack?

- ► .NET / C# / Microsoft ecosystem
- ► Java / JVM (Spring, Kotlin)
- ► Python (data science, ML, backend)
- ► Node.js / JavaScript / TypeScript
- ► Go / Rust / C++
- ► Mixed / polyglot

Weight: Medium **Branching note:** *.NET strongly boosts Azure; Python boosts GCP; mixed is neutral.*

Q19 What CI/CD toolchain does your team use?

- ► Azure DevOps / GitHub Actions
- ► GitLab CI
- ► Jenkins / CircleCI / other self-hosted
- ► No CI/CD pipeline yet

Weight: Low **Branching note:** *Azure DevOps gives slight Azure preference; otherwise low impact.*

Q20 Do you currently use any vendor-specific on-premise products you wish to migrate?

- ► Microsoft (Active Directory, SQL Server, Exchange)
- ► Oracle (Oracle DB, Oracle ERP / Fusion)

- IBM (Db2, WebSphere, IBM MQ, AIX)
- VMware / Dell
- None — open-source or neutral stack
- Multiple vendors

Weight: HIGH (×2.5) **Branching note:** *Existing vendor products are the single strongest predictor of provider fit.*

Q21 How important are managed AI and machine learning services to your product roadmap?

- Critical — AI is core to our product
- Important — we plan AI features within 12 months
- Nice to have — possible future consideration
- Not relevant

Weight: Medium **Branching note:** *Critical/Important answers boost GCP (Vertex AI, TPUs) and AWS (SageMaker).*

Stage E — Budget, Operations & Timeline (Q22–Q25)

Q22 What is your estimated monthly cloud infrastructure budget?

- Under $500 / month
- $500 – $5,000 / month
- $5,000 – $50,000 / month
- $50,000 – $500,000 / month
- Over $500,000 / month

Weight: High **Branching note:** *Budget level determines discount eligibility and viable support tiers.*

Q23 Do you need to maintain connectivity to on-premise systems after migration?

- Yes — ongoing hybrid connectivity (VPN or dedicated connection)
- Yes — temporary bridge during migration only
- No — full cloud migration, all systems moving

Weight: HIGH (×2.5) **Branching note:** *Hybrid requirement significantly boosts Azure (Azure Arc) and AWS (Outposts).*

Q24 What level of cloud support do you require?

- Community only (forums, documentation)
- Business support (email / chat response within hours)
- Enterprise support (24/7 phone, dedicated TAM)
- Mission-critical (sub-15-minute response SLA)

Weight: Medium **Branching note:** *Mission-critical support boosts AWS and Azure which have the most mature enterprise support programs.*

Q25 What is your target migration timeline?

- ► Immediate — within 1–3 months
- ► Short-term — 3–6 months
- ► Medium-term — 6–12 months
- ► Long-term — 12–24 months
- ► Not yet defined

Weight: Medium **Branching note:** *Aggressive timelines favour providers with superior migration tooling and free migration credits.*

5. User Experience Flow Design

The UX follows a 5-screen wizard. The single-question-per-screen format from the ChatGPT draft is adopted in full. Real-time scoring previews appear after Stage C (compliance questions) to build engagement.

Screen 1 — Welcome / Introduction

Headline: Find the Perfect Cloud Provider for Your Infrastructure
Sub-copy: Answer 12–20 questions tailored to your setup. Takes about 5–8 minutes.
CTA button: Start Assessment
Trust signal: 'Used by 10,000+ teams planning their cloud migration'

Screen 2 — Staged Progress Indicator

Persistent progress bar at top: Stage 2 of 5 — Application Architecture
Question counter: Question 6 of ~18 (updates dynamically as branching skips questions)
Back button available on every screen — answers are preserved
Auto-save every question so users can resume after closing the browser

Screen 3 — Question Interface — Visual Answer Cards

One question displayed per screen in large, readable typography
Answers presented as selectable cards with icon + label + one-line description
Example — Workload Type cards:
[Web Applications] Host websites, APIs, and customer-facing services
[AI / ML] Train and deploy machine learning models

[Data Analytics] Large-scale data pipelines and reporting
[Enterprise Apps] ERP, CRM, and internal business applications
Multi-select enabled where noted (e.g., compliance frameworks, workload types)
Tooltip '?' icon on every technical term links to a plain-English explanation

Screen 4 — Real-Time Recommendation Preview (After Stage C)

Shown after compliance questions are answered — when scores are most differentiated
Live animated bar chart showing running scores:

Subtitle: 'Complete the remaining questions to refine your recommendation'
Purpose: increases completion motivation by making value visible early

Screen 5 — Final Recommendation Dashboard

Recommended Provider badge with provider logo and confidence score (e.g., 78 / 100)
Why [Provider] fits your needs — 3 bullet points generated from answer evidence
Trade-off callout — 1–2 honest considerations before committing
Alternative Provider — second-ranked provider with one-line differentiation note
Migration Complexity Estimate — Low / Medium / High with time range
Next Steps panel — links to architecture generator, migration plan, cost estimator
Download PDF Summary button — exports answers + recommendation + rationale

5.1 Smart Branching Rules

Condition	Action
No machine learning workload selected	Skip Q21 (AI/ML services importance)
Oracle DB selected in Q8	Skip other DB scoring; apply Oracle Cloud hard-boost
No compliance framework selected in Q14	Skip Q16 (data sensitivity) for non-regulated users
No DR requirement in Q17	Skip multi-region architecture questions
Full cloud migration in Q23	Skip hybrid connectivity questions
No cloud experience in expertise question	Surface simplified provider summaries on result screen

6. Recommendation Explanation Logic

Every recommendation must be explainable. The engine generates three components to build user trust and transparency.

6.1 Evidence Block

Identifies the 3–4 user answers that most influenced the recommendation. Example for an AWS result:

- Your workload is primarily web applications and APIs → AWS leads this category
- You selected an open-source / Linux stack → AWS is the strongest fit
- You require rapid auto-scaling → AWS Auto Scaling is the most mature option
- Your monthly budget is $5,000–$50,000 → AWS Reserved Instances offer strong savings at this tier

6.2 Comparison Block

Shows where the recommended provider leads vs. alternatives on the user's specific priorities:

Your Priority	AWS	GCP	Azure
Scalability	★★★★★	★★★★☆	★★★★☆
Developer Ecosystem	★★★★★	★★★★☆	★★★★☆
Migration Tooling	★★★★★	★★★☆☆	★★★★★
Cost Efficiency	★★★☆☆	★★★★★	★★★☆☆

6.3 Transparency & Trade-offs Block

Honest caveats are displayed below the recommendation to prevent blind commitment. Example for AWS:

- AWS pricing is more complex than GCP — budget for a billing review in your first 90 days
- Egress fees can be significant if you pull large datasets back out of S3 regularly

6.4 Migration Complexity Estimate

How calculated	Complexity = f(data volume, hybrid requirement, compliance workload, existing vendor lock-in, team cloud experience)

Level	Criteria	Typical Timeline
Low	< 1 TB data, no compliance, stateless apps, experienced team	1–2 weeks

Medium	1–100 TB, basic compliance, some stateful services, mixed team	2–8 weeks
High	> 100 TB, regulated industry, hybrid required, legacy systems	2–6 months

7. Technical Implementation Requirements

7.1 Frontend

- Multi-step wizard with persistent progress bar (Stage X of 5, Question Y of ~N)
- Answer type per question: single-select card, multi-select card, or slider as specified
- Conditional branching evaluated client-side for instant response — no round-trip to server
- Inline tooltip '?' on every technical term; tooltip text stored in a JSON config file
- Auto-save to sessionStorage every question answer; full restore on page refresh
- Mobile-responsive: minimum 375px viewport, card layout stacks vertically on mobile
- Animated score preview chart (horizontal bar) updates after each Stage C question

7.2 Scoring Engine (Backend)

- Stateless pure function: input = JSON answer payload, output = ranked provider list with scores and rationale strings
- Scoring deltas and category weights stored in YAML config — tunable without code deploy
- Hard-eliminator rules run first, before weighted scoring, as a filter pass
- Rationale strings generated from a deterministic template engine (not LLM) for consistency
- Answer payload logged (anonymized, no PII) to data warehouse for future model refinement
- P95 response time target: < 200ms for scoring computation

7.3 Answer Payload Schema (JSON)

```
{ "sessionId": "uuid-v4",
  "completedAt": "2026-03-08T10:30:00Z",
  "answers": {
    "q1_data_volume":        "1tb_to_10tb",
    "q5_multiregion":        "active_passive",
    "q7_containers":         "yes_production",
    "q8_database_type":      ["relational", "nosql"],
    "q14_compliance":        ["hipaa", "soc2"],
    "q15_data_residency":    "country_specific",
```

```
    "q20_existing_vendor":   "microsoft",
    "q22_monthly_budget":    "5k_to_50k",
    "q23_hybrid_required":   "yes_ongoing"
  }
}
```

8. Acceptance Criteria & Success Metrics

8.1 Acceptance Criteria

1. Full questionnaire completable in under 8 minutes without prior cloud knowledge
2. All 25 questions render correctly on desktop (≥1024px) and mobile (≥375px) with branching
3. Hard-eliminator rules exclude incompatible providers correctly in 100% of defined test cases
4. Scoring engine returns recommendation within 200ms (P95) of questionnaire submission
5. Each recommendation includes a minimum of 3 evidence bullets and 2 trade-off callouts
6. PDF export renders with user answers, final scores, rationale, and migration estimate
7. Session state persists across browser refresh for up to 24 hours via sessionStorage
8. Real-time score preview chart visible and updating after Stage C questions

8.2 Success KPIs (Post-Launch)

KPI Metric	Target — 30 days	Target — 90 days
Questionnaire completion rate	> 70%	> 80%
Avg. session duration	< 8 min	< 6 min
User satisfaction (1–5 survey)	> 4.0	> 4.3
Recommendation acceptance rate	> 50%	> 60%
Report / PDF download rate	> 30% of completions	40%
Support tickets re: cloud choice	Baseline set	< 20% of baseline

9. Future Roadmap

Phase 2 — Live Cost Estimation

- Integrate live pricing APIs from AWS, Azure, GCP, and OCI to generate monthly cost breakdowns based on workload parameters from the questionnaire
- Add a 'What-if' comparison mode: run side-by-side cost projections with different configuration assumptions
- Display a migration cost estimate (one-time engineering effort + ongoing infrastructure cost)

Phase 3 — Guided Migration Playbook

- Generate a step-by-step migration checklist tailored to the selected provider and specific workload type
- AI-powered chat assistant for follow-up questions after the recommendation is generated
- Automated architecture diagram suggestions based on workload and provider answers
- Multi-cloud recommendation mode: distribute workloads intelligently across providers
- Integration with cloud provider partner programs to connect users with certified migration partners

Version	Authors	Status	Date
v2.0 (Combined)	ChatGPT + Claude	Draft — For Review	March 2026